

---

# PULER

---

Bassel El Mabsout

April, 2023

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Compiler design</b>	<b>2</b>
2.1	Recursion-Schemes . . . . .	3
2.2	Other compiler stages . . . . .	4
2.2.1	Renamer . . . . .	4
2.2.2	Interpreter . . . . .	5
2.2.3	Autoformatter . . . . .	5
2.2.4	Emission . . . . .	5
2.3	REPL . . . . .	6
<b>3</b>	<b>PULER</b>	<b>6</b>
3.1	Syntax . . . . .	6
3.2	Features . . . . .	7
3.3	Type system . . . . .	7
<b>4</b>	<b>Examples</b>	<b>9</b>
4.1	Example Program . . . . .	9
4.2	Example Commands . . . . .	9
4.2.1	REPL example . . . . .	9
4.2.2	Compilation stages example . . . . .	10
<b>5</b>	<b>Installation</b>	<b>11</b>
<b>6</b>	<b>Conclusion</b>	<b>12</b>

# 1 Introduction

Typical programming language compilers are comprised of various *stages* that are chained together producing a final output from an initial input program. These stages can include parsing, renaming, type inference, optimization, and emission. For example, a C compiler takes as input a C program, which goes through these transformations and finally produces machine code. Similarly the Purescript [7] compiler produces JavaScript. While these transformations share abstract language constructs like **let-in** statements, **lambda** abstractions, and **if-then-else** statements [19], they also require transformation-specific information to be tracked. For instance, during type inference, a data-structure storing the language's expressions **decorated** with their types and ununified variables must be defined. [2]

When compilers are implemented in a programming language supporting **Abstract Data Types (ADTs)** [9], they often contain multiple ADTs for each transformation, leading to duplication of effort, functionality, and difficulty in keeping the language's *constructs* in sync. Additionally, extending the language with new constructs becomes difficult as the language designer must add necessary mechanisms to every datatype and function. Without a core datatype, bugs may arise since some intermediary ADTs may be left untouched, and the compiler may not reject the updated code. This issue presents itself in multiple different compiler implementations [14, 1], and is the core focus of a body of existing work [24, 3, 13].

Motivated by the ADT duplication problem, we explore another point in the *design space* by introducing an experimental compiler with a differing architecture, and studying its implications in Section 2. This compiler accepts PULER programs, a new experimental *ML-based* programming language presented in Section 3 of this document. It then invokes different transformation as required by the final output requested, for an example of the stages involved, see Section 4.2.2 One of the distinctive features of PULER is defining **unification rules** for type mismatches, forming types themselves and being treated as *first-class* citizens in the type system. In contrast, existing compilers mostly exit and produce an error upon a type mismatch. Existing works focusing on this issue of type error debugging [18] generate orthogonal constraints which are separately represented and solved (e.g. via **SMT** solvers), producing errors which help users figure out the root cause. These ideas are discussed in Section 3.3.

# 2 Compiler design

On the implementation of compilers written in Haskell specifically, the work by Najd and Peyton Jones [17] presents and implements "Trees that grow", a technique for "decorating" ADTs by using type families. They then rewrite the internals of the Glasgow Haskell Compiler (GHC) [14] to make use of the method. In GHC, the data types for these transformations are large (dozens of types, hundreds of constructors) making them difficult to maintain. Their solution involves defining a single extensible ADT with Haskell's possible expressions. This allows different "decorations" of Haskell's Abstract Syntax Tree (AST) to be defined in separate parts of the compiler. Plugin writers then benefit from reusing the compiler as a library when extending the language.

In the PULER compiler, each language construct is given a separate core parameterized type which is then reused across transformations. We then make use of deriving mechanisms to create reusable functionality across trees composed of these types. For example the following is the definition for the datatype of a Let expression:

```
data Let expr = Let (Decs expr) expr
  deriving (Eq, Functor, Foldable, Traversable)
```

This technique allows us to define dependencies between language constructs (here let expressions use declarations) without constraining the final tree structure and decorations added. The language designer can then work with each structure on its own, and then define an ADT combining them. This allows the separation of the logic handling each type of expression in the language from the logic handling their composition. As an example, we show how the parser for let expressions is defined:

```
instance Parsable a => Parsable (Let a) where
  parser = Let
    <$> (begKeyword "let" *> parser)
    <*> (midKeyword "in" *> parser)
```

Typeclasses allow us to define how each piece of the language can be parsed while remaining generic with respect to how subexpressions are composed. Contrasting this with the "Trees that grow" method, the structure of the language is solidified in a single ADT, and while they allow open extension to some parts of the language, they cannot introduce new language constructs. More information about the syntax of PULER is provided in Section 3.1.

## 2.1 Recursion-Schemes

In order to recursively connect subexpressions together, we form a core ADT:

```
data Expr
  = Evar Var
  | Eapp (App Expr)
  | Elet (Let Expr)
  | Eif (If Expr)
  | Edec (Decs Expr)
  | Elambda (Lambda (N.NonEmpty Var) Expr)
  | Efix (Fixer Expr)
  | Elit Lit
  | Eannotation (Annotated Expr NamedTypes)
  deriving (Eq, Generic)
makeBaseFunctor ''Expr
```

Recursion-schemes [16] allow recursive algorithms to decouple local computations from the patterns used in the recursive step. This allows us to treat our **Expr** datatype as an *F-algebra* [4] and write our transformations as combinations of algebras and co-algebras over similar types. Specifically, we make use of Haskell's recursion-schemes [11] library. Combined with the use of effect systems (Polysemy [21]), this allowed breaking down the transformations into many smaller subproblems, and then composing them together into the final computed value. The "makeBaseFunctor" template function creates an auxiliary datatype, namely "ExprF" which now acts as our interface with recursion schemes. This general method of structuring the compiler, allowed us to define different ADTs specific to each transformation, yet shared information is retained via the different datatypes for each construct. For example, here's the definition of the ADT representing the possible values that the interpreter can produce:

```

data Value = Vlit Lit
           | Vlambda (Lambda (N.NonEmpty Var) LambdaBody)
           | VDecs (Decs Value)
           | Empty

```

Notice how, even though we define a new ADT that must be changed when adding new language constructs, the actual datatypes are just wrappers connecting our general datatypes (such as Decs). This trades off some duplication for flexibility, as such this occupies a point in the design space between full duplication and "Trees that grow".

## 2.2 Other compiler stages

This section contains the details of other compilation stages and of their implementation. They all make use of the aforementioned ideas for making compiler structuring decisions.

### 2.2.1 Renamer

The renamer keeps track of the scope of each variable and renames overlapping variables to an incremented version so that we don't worry about scope in later compiler passes:

```

-----INPUT-----
x = 3;
x = 4;
z = "test";
y = \x -> let f = 4 in f + x;
main = x;
-----RENAMED-----
x = 3;
x#1 = 4;
z = "test";
y = \x#2 -> let f = 4;
           in ((Add f) x#2);
main = x#1;

```

This is the main renaming logic which just allows us to push and pop a variable stack:

```

data Scope = Scope {numShadowing :: Shadowing, stack :: [Shadowing]}

type Renamer = M.Map Var Scope

insertToScope :: Var -> Renamer -> Renamer
insertToScope =
  M.alter $ Just . \case
    Just (Scope g l) -> Scope (g + 1) (g + 1 : l)
    Nothing -> Scope 0 $ [0]

removeFromScope :: Var -> Renamer -> Renamer
removeFromScope = M.update \ (Scope g l) -> Just (Scope g (tail l))

```

A catamorphism is then used to recursively update the variables associated with individual scope levels, producing the final transformed AST.

### 2.2.2 Interpreter

This is PULER's interpreter, it works by first "blinding" the original expression tree, meaning it hides the body of functions from recursion schemes so that we never evaluate the values inside of the bodies of functions. It then converts the AST into a Value type (as defined above). This is done with a catamorphism with access to a monad that keeps track of the scope of variables. The interpreter even correctly evaluates the program without a renaming step.

### 2.2.3 Autoformatter

This is the autoformatter that comes with PULER, it appropriately indents new lines if there isn't too much width space. Here's a short program:

```
x = \test -> if True then 1 else 2
-----PARSED-----
x = \test -> if True
           then 1
           else 2;
```

Here's a longer one:

```
-----INPUT-----
x = \test -> if True then "this is a very long message" else let y = 3 in ""
-----PARSED-----
x = \test ->
    if True
    then "this is a very long message"
    else let y = 3;
        in "";
```

Notice that it decided to insert a new line after  $\rightarrow$  in the second case so that it can fit within 40 characters. Similar to how the parser is defined, each datatype has its own pretty printer, which is then stringed together forming the full pretty-printable expression type.

### 2.2.4 Emission

Finally, we emit the typed language to Python. The output is auto-formatted so it retains some readability even after transformation. The runtime is fairly simple, consisting of a few functions allowing the direct translation of any PULER program into python. There are some extra tools in the emission code allowing the computation of the free variables computing the environment of every function and such, but for python since we can create lambdas, then there's no need to create closures ourselves. The compiler initially was targeting C, however the combination of partial applications and C having no ability to create lambdas meant, if done naively, the generated output would create n functions for every n variable function which was very unappealing. Another possibility is Mallocing the environment of a function upon creation. Storing it in a struct that also stores the arguments to the function, then when applying a single argument, simply assign the value to the argument in the struct. Now our function would be passed as a struct of arguments and an environment. When the last argument is provided, that's where we actually call the function. The only problem with this approach is that it requires tagging the functions to figure out when the

final application is happening. The only ways of doing this that I've thought of end up coupling the language to the backend too much. So I opted to emit to a language with closures. Namely, python.

## 2.3 REPL

The language also includes a Read Eval Print Loop, which allows users to run valid (or invalid) PULER and interact with the language. This works by using the context associated to every stateful operation in our architecture. Meaning the same code that tracks information like the current scope of variables is reused in the repl to keep this state live. We can also get the benefit of providing features like autocomplete since the context is mostly tracked using dictionaries. Another feature is that we can load files into the repl, if you write a PULER program and save it, you can start a repl that evaluates it and allows you to write code with it as part of the context. I found this to be immensely useful when debugging.

## 3 PULER

PULER is an experimental ML based programming language that is strictly evaluated, partially applied, implements a Hindley-Milner-based [5] type system (Section 3.3), and is lexically scoped (Section 2.2.1). It can be interpreted (Section 2.2.2) or compiled (Section 2). There's also a Read Eval Print Loop (REPL) (Section 2.3) available so that one can interact with the language live. The core language is simple, containing a few data types and language constructs. It can be extended through either adding compiler stages or minimally changing the representation used by the compiler to introduce new features. The language is functional first, meaning functions are first class citizens, and it disallows mutation, making it a referentially-transparent language.

One interesting feature of PULER is that the type system represents type mismatches. Unification rules are defined for mismatching types. This contrasts with typical programming languages' type checking implementations, which exit and produce an error immediately upon mismatch. We motivate

### 3.1 Syntax

PULER's syntax is very similar to languages like SML and Haskell. The following is a high-level BNF[15] with the details in Parser.hs:

$$\begin{aligned}
\langle abs \rangle &::= \backslash \langle var \rangle \rightarrow \langle subexpr \rangle \\
\langle subexprs \rangle &::= \langle subexpr \rangle | \langle subexpr \rangle \langle subexprs \rangle \\
\langle app \rangle &::= (\langle subexpr \rangle \langle subexprs \rangle) \\
\langle lit \rangle &::= \langle string \rangle | \langle integer \rangle | \langle boolean \rangle | \{\} \\
\langle dec \rangle &::= \langle variable \rangle = \langle subexpr \rangle \\
\langle decs \rangle &::= \langle dec \rangle | \langle dec \rangle ; \langle decs \rangle \\
\langle let \rangle &::= \text{let } \langle declarations \rangle \text{ in } \langle subexpr \rangle \\
\langle fix \rangle &::= \text{fix } \langle var \rangle \langle abs \rangle \\
\langle type \rangle &::= \langle type \rangle \rightarrow \langle type \rangle | \{\} | Int | Bool | Str \\
\langle annotation \rangle &::= \langle subexpr \rangle : \langle type \rangle \\
\langle subexpr \rangle &::= \langle abs \rangle | \langle lit \rangle | \langle let \rangle | \langle fix \rangle | \langle annotation \rangle | \langle app \rangle
\end{aligned}$$

At the top level of a PULER file  $\langle decs \rangle$  are expected.

## 3.2 Features

- Lexical scoping: The scope of every variable is exactly defined by the location of its definition. Meaning its scope is completely defined at compile time. Defining a variable with the same name as another variable that is in scope means that this new variable shadows the earlier defined one.
- Partially applied: Applying an argument to a 3 argument function returns a function of 2 arguments instead of being an error.
- Hindley-Milner-based type system: If the program is correct then no types will need to be written anywhere, it is all inferred.
- Strict evaluation: The language is strictly evaluated, meaning function arguments are evaluated before the function is.
- Immutable: Any operation or function cannot modify the value of an existing variable. This along with the ability to create pure functions makes the language referentially transparent (this behavior is not preserved when you use the only impure function in PULER which is `Print`). Meaning you can replace a variable by the statement used to create it. The `=` operator is more like the mathematical definition of `=` and less like the assignment we're used to in imperative programming languages.

## 3.3 Type system

Multiple previous works have focused on the problem of type error localization. Some show relevant portions of failed type inference traces [6, 25], others show a slice of the program involved in the error [8, 23]. Other methods repeatedly call the typechecker, finding several error sources [12]. The most promising approach we've observed in this area comes from the works Pavlinovic et al. [18], Stuckey et al. [22], producing constraints which rank type errors and then are solved by an SMT solver. PULER makes no such choices, instead of producing an error and exiting the computation, the PULER type system keeps track of

type mismatches. This is orthogonal to the works by Pavlinovic et al. [18] as then the decision to show some errors can be taken later, or the user can be shown multiple different “views” of the possible errors that happened for flexibility in type error debugging. Note that the PULER compiler cannot currently handle all kinds of type errors this way, as an infinite occurs check (for preventing infinite types) also results in the compiler exiting.

Type inference in PULER converts an `Expr` to a `Cofree ExprF INamedTypes`, this just means that we’re annotating each node of the recursive abstract syntax tree with a type. The type inference algorithm does the following: It initializes each expression with unification variables and relates the unification variables with inference rules. Then we “propagate” [20] the knowledge gained through a pass of applying the inference rules. We repeat this process (while doing occurs checks) until no further knowledge can be gained. When this is done we return the annotated tree. This is the datatype describing the types used for inference:

```
data Itypes base
  = Iunif Name
  | IbaseType base
  | Iarrow (Itypes base) (Itypes base)
  | Imismatch [Itypes base]
```

In order to precisely describe what it means for knowledge to increase, we allow any 2 types to be joined to create a type considered the least upperbound of both:

```
instance Ord base => Semigroup (Itypes base) where
  -- this is actually a join semilattice
  a <> b | a == b = a
  Imismatch as <> Imismatch bs = Imismatch (nubQuick (as ++ bs))
  Iarrow a1 b1 <> Iarrow a2 b2 = Iarrow (a1 <> a2) (b1 <> b2)
  a@(Iunif _) <> (Iunif _) = a
  (Iunif a) <> b = b
  a <> (Iunif b) = a
  (Imismatch l) <> x = if x `elem` l then x else Imismatch (nubQuick $ x:l)
  x <> (Imismatch l) = if x `elem` l then x else Imismatch (nubQuick $ x:l)
  a <> b = Imismatch [a, b]
```

When we have 2 different types we can now join them to create a “larger” type. We now proceed by unifying types that are supposed to be equal (because of rules on the structure of the expression tree) by representing the types that are equal to each other as a disjoint set. Then we find the least upper bound of each disjoint set. Knowledge is then described as a map between each type and its corresponding “joined” type. Once we have this map we can “gain” knowledge by replacing any unification variables with a corresponding concrete type. This in turn will generate new types to unify. This process is repeated until this knowledge map is unchanging. What’s nice about doing things this way is that we can give users as much information about the types of expressions as we can get even when the types mismatch or are kept ambiguous. Here’s an example with mismatching and ununified types:

```
x = 4;
y = x + "string";
main = (x 3);
-----INFERRED-----
```



```

x = 4:Int;
y = ((Add x:[ Int
           , String
           , Int->?d ]):[ Int
                        , String ]->Int
     "string":String):[ Int
                       , String
                       , Int->?d ];
main = (x:[Int, String, Int->?d]
        3:Int):?d;:{}

```

Notice that there are 3 different type errors in here. First when we see the type `[a, b]` this means that the types `a` and `b` don't match. We can see here that the type of `x` is inferred to be `[Int, String, Int->?d]`. Second when we see a `?` symbol this means it is a unification variable and that this variable remained ununifiable. Then notice that `x` it's self is actually used in three separate incompatible ways. First `x` is set to be the number 4 which is an `Int`, then we're adding a string to `x` but you can't add strings to numbers so that creates the first mismatch, the second mismatch comes from using `x` as a function in `main`. Given that we never use the result of the function this means we don't "know" what the result type is, which is another type error. Even though the output is currently verbose we could simply query the type of specific variables and this ability allows tracking the trace of a mismatch which we also found useful.

## 4 Examples

### 4.1 Example Program

```

fib = fix fib \x ->
  if x == 0
  then 0
  else if x == 1
  then 1
  else (fib (x-1)) + (fib (x-2));

main = (fib 8);
>>
21

```

### 4.2 Example Commands

#### 4.2.1 REPL example

Open a repl and run expressions as well as check their types:

```

> PULER repl
Welcome to PULER!

> x = 1
x = 1;

```

```

> y = \z -> z
y = \z -> <body>;
> :t (y x)
(y: Int -> Int x: Int): Int
> (y x) + 1
2
> (Print "test")
"test"
{}

```

#### 4.2.2 Compilation stages example

Compile an example \*.pul file implementing the factorial function

```
> PULER example.pul
```

```

-----INPUT-----
fact = fix f \x ->
    if x == 0
    then 1
    else let nextFac = x * (f (x - 1));
          p = (Print (Int2Str nextFac))
          in nextFac;
main = let g = (fact 8) in {};

-----PARSED-----
fact = fix f \x -> if ((EqInt x) 0)
    then 1
    else let nextFac = ((Mul x) (f ((Sub x) 1)));
          p = (Print (Int2Str nextFac));
          in nextFac;
main = let g = (fact 8);
    in {};

-----RENAMED-----
fact = fix f \x -> if ((EqInt x) 0)
    then 1
    else let nextFac = ((Mul x) (f ((Sub x) 1)));
          p = (Print (Int2Str nextFac));
          in nextFac;
main = let g = (fact 8);
    in {};

-----INFERRED-----
fact = fix f \x -> if ((EqInt x:Int):Int->Bool 0:Int):Bool
    then 1:Int
    else let nextFac = ((Mul x:Int):Int->Int
        (f:Int->Int ((Sub x:Int):Int->Int
            1:Int):Int):Int;
          p = (Print (Int2Str nextFac:Int):String):{};
          in nextFac:Int:Int:Int:Int->Int;

```

```

main = let g = (fact:Int->Int 8:Int):Int;
      in {}:{}:{};:{}
-----CHECKED-----
Typechecks!
-----EVALUATED-----
1
2
6
24
120
720
5040
40320
{}
-----EMITTED-----
fact = fix(lambda f: lambda x:(
  (
    1
  ) if (
    EqInt(x)(0)
  ) else (
    let(
      nextFac = Mul(x)(f(Sub(x)(1))),
      inn = lambda nextFac:(
        let(
          p = Print(Int2Str(nextFac)),
          inn = lambda p:(
            nextFac
          )
        )
      )
    )
  )
))
main = let(
  g = fact(8),
  inn = lambda g:(
    unit
  )
)

```

## 5 Installation

Install the compiler by following these instructions:

1. Clone the following repository: [PULER](#)
2. Install the [nix package manager](#)
3. Navigate to the PULER folder
4. Run the `nix-shell` command which puts you in a reproducible environment with the exact packages (and packages versions) required to compile the compiler

5. Compile the code by running `ghc -O2 Main.hs`
6. Now you can use the Compiler by calling `./Main`

## 6 Conclusion

After multiple iterations of compiler design, PULER has converged to the idea of having main datatypes stringed together to form algebras on which the algorithms forming compiler stages are built upon. In exploring this architecture, we’ve built a full programming language with multiple interesting features. The architecture made it easy to use the compiler as a library and build a REPL for example with very minimal amounts of code. We did notice however, that even with careful design in avoiding duplication, there is still boilerplate code which duplicates parts of the core ADT throughout compiler transformations. The sum types used tag each alternative branch with a name which is unnecessary in our approach since the datatypes are already unique. As such, we believe connecting the types of expressions together via strongly typed heterogeneous lists [10] would lead to an even more terse and composable interface, we leave this as an exploration in future work. At its current state, we found it easier to add any type of construct to the language and fill in the implementation gaps, than otherwise would be the case without this architecture. Another remaining challenge is targetting C as a backend, the limitations of which may strain the architectural decisions made in PULER.

## References

- [1] Ocaml. URL <https://v2.ocaml.org/manual/>.
- [2] A. W. Appel and J. Palsberg. Introduction. In *Modern Compiler Implementation in Java*, pages 3–15. Cambridge University Press, Cambridge, Oct. 2002.
- [3] E. Axelsson. A generic abstract syntax model for embedded languages. In *Proceedings of the 17th ACM SIGPLAN International Conference on Functional Programming, ICFP ’12*, page 323–334, New York, NY, USA, 2012. Association for Computing Machinery. ISBN 9781450310543. doi: 10.1145/2364527.2364573. URL <https://doi.org/10.1145/2364527.2364573>.
- [4] M. Bartosz. F-algebras. URL <https://bartoszmilewski.com/2017/02/28/f-algebras/>.
- [5] L. Damas and R. Milner. Principal type-schemes for functional programs. In *Proceedings of the 9th ACM SIGPLAN-SIGACT symposium on Principles of programming languages - POPL 82*. ACM Press, 1982. doi: 10.1145/582153.582176. URL <https://doi.org/10.1145/582153.582176>.
- [6] D. Duggan and F. Bent. Explaining type inference. *Science of Computer Programming*, 27(1):37–83, 1996. ISSN 0167-6423. doi: [https://doi.org/10.1016/0167-6423\(95\)00007-0](https://doi.org/10.1016/0167-6423(95)00007-0). URL <https://www.sciencedirect.com/science/article/pii/0167642395000070>.
- [7] P. Freeman. *PureScript by Example*. <https://leanpub.com/purescript>, 2017. <https://leanpub.com/purescript>.

- [8] H. Gast. Explaining ML type errors by data flows. In *Implementation and Application of Functional Languages*, pages 72–89. Springer Berlin Heidelberg, 2005. doi: 10.1007/11431664\_5. URL [https://doi.org/10.1007/11431664\\_5](https://doi.org/10.1007/11431664_5).
- [9] J. Gutttag and J. Horning. The algebraic specification of abstract data types. *Acta Informatica*, 10(1), 1978. doi: 10.1007/bf00260922. URL <https://doi.org/10.1007/bf00260922>.
- [10] O. Kiselyov, R. Lämmel, and K. Schupke. Strongly typed heterogeneous collections. In *Proceedings of the 2004 ACM SIGPLAN Workshop on Haskell*, Haskell ’04, page 96–107, New York, NY, USA, 2004. Association for Computing Machinery. ISBN 1581138504. doi: 10.1145/1017472.1017488. URL <https://doi.org/10.1145/1017472.1017488>.
- [11] E. Kmett. Recursion-schemes. URL <https://hackage.haskell.org/package/recursion-schemes>.
- [12] B. S. Lerner, M. Flower, D. Grossman, and C. Chambers. Searching for type-error messages. *SIGPLAN Not.*, 42(6):425–434, jun 2007. ISSN 0362-1340. doi: 10.1145/1273442.1250783. URL <https://doi.org/10.1145/1273442.1250783>.
- [13] A. Löb and R. Hinze. Open data types and open functions. In *Proceedings of the 8th ACM SIGPLAN International Conference on Principles and Practice of Declarative Programming*, PPDP ’06, page 133–144, New York, NY, USA, 2006. Association for Computing Machinery. ISBN 1595933883. doi: 10.1145/1140335.1140352. URL <https://doi.org/10.1145/1140335.1140352>.
- [14] S. Marlow and S. L. P. Jones. The glasgow haskell compiler. 2012.
- [15] D. D. McCracken and E. D. Reilly. *Backus-Naur Form (BNF)*, page 129–131. John Wiley and Sons Ltd., GBR, 2003. ISBN 0470864125.
- [16] E. Meijer, M. Fokkinga, and R. Paterson. Functional programming with bananas, lenses, envelopes and barbed wire. In *Functional Programming Languages and Computer Architecture*, pages 124–144. Springer Berlin Heidelberg, 1991. doi: 10.1007/3540543961\_7. URL [https://doi.org/10.1007/3540543961\\_7](https://doi.org/10.1007/3540543961_7).
- [17] S. Najd and S. Peyton Jones. Trees that grow. *Journal of Universal Computer Science (JUCS)*, 23:47–62, January 2017. URL <https://www.microsoft.com/en-us/research/publication/trees-that-grow/>.
- [18] Z. Pavlinovic, T. King, and T. Wies. Finding minimum type error sources. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages and Applications*, OOPSLA ’14, page 525–542, New York, NY, USA, 2014. Association for Computing Machinery. ISBN 9781450325851. doi: 10.1145/2660193.2660230. URL <https://doi.org/10.1145/2660193.2660230>.
- [19] B. C. Pierce. *Types and Programming Languages*. The MIT Press. MIT Press, London, England, Jan. 2002.
- [20] A. Radul. *Propagation networks: A flexible and expressive substrate for computation*. PhD thesis, Massachusetts Institute of Technology, 2009.

- [21] M. Sandy. Polysemy. URL <https://hackage.haskell.org/package/polysemy>.
- [22] P. J. Stuckey, M. Sulzmann, and J. Wazny. Improving type error diagnosis. In *Proceedings of the 2004 ACM SIGPLAN Workshop on Haskell*, Haskell '04, page 80–91, New York, NY, USA, 2004. Association for Computing Machinery. ISBN 1581138504. doi: 10.1145/1017472.1017486. URL <https://doi.org/10.1145/1017472.1017486>.
- [23] F. Tip and T. B. Dinesh. A slicing-based approach for locating type errors. *ACM Trans. Softw. Eng. Methodol.*, 10(1):5–55, jan 2001. ISSN 1049-331X. doi: 10.1145/366378.366379. URL <https://doi.org/10.1145/366378.366379>.
- [24] M. Torgersen. The expression problem revisited. In *ECOOP 2004 – Object-Oriented Programming*, pages 123–146. Springer Berlin Heidelberg, 2004. doi: 10.1007/978-3-540-24851-4\_6. URL [https://doi.org/10.1007/978-3-540-24851-4\\_6](https://doi.org/10.1007/978-3-540-24851-4_6).
- [25] M. Wand. Finding the source of type errors. In *Proceedings of the 13th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages*, POPL '86, page 38–43, New York, NY, USA, 1986. Association for Computing Machinery. ISBN 9781450373470. doi: 10.1145/512644.512648. URL <https://doi.org/10.1145/512644.512648>.